

CORTEX XSIAM • CONTENT PACK

KOI

Endpoint and agentic-AI security — collected, investigated and acted on, automatically.

26

commands

10

playbooks

1

dashboard

2

rule sets

THE PROBLEM

The software you never deployed — but still run

KOI discovers, scores and governs the long tail of software your workforce installs itself, including the fast-growing agentic-AI surface.

1

Browser extensions

Chrome, Edge and friends — the classic blind spot.

2

IDE plugins

VS Code and JetBrains add-ons running with developer rights.

3

Code packages

npm and PyPI dependencies pulled straight into builds.

4

Desktop applications

Installed outside any deployment pipeline.

5

MCP servers

Model Context Protocol endpoints wired into AI agents.

6

AI models & skills

The rest of the agentic-AI supply chain.

CONTENTS

What ships in the pack

A

KOI integration

Event collector plus 26 commands across the KOI API.

B

Parsing & modeling rules

Normalize raw events and map them to the Cortex Data Model.

C

Alerts dashboard

Ready-made XSIAM dashboard for KOI alert activity.

D

Triage & response playbooks

Seven playbooks: triage, investigation, and gated response.

E

Script Runner playbooks

Three playbooks to run KOI scripts on Cortex-Agent endpoints.

Two independent automation tracks

Detect & respond

Each KOI alert is triaged, investigated and routed — with response gated on an analyst.

Seven playbooks: triage, two investigations, enrichment, response and hunting.

Operate at fleet scale

Run KOI script packages on Cortex-Agent endpoints on a schedule, targeted by OS, group or hostname.

Three playbooks, configured entirely through a JSON List.

Everything installs together as one pack.

What you can do with it

1

Telemetry collection

Alerts and audit logs into koi_koi_raw, mapped to XDM.

2

Asset & posture visibility

Everything running, per item or per device, with risk scoring.

3

Threat intelligence

Koidex risk score, findings and an AI-generated risk summary.

4

Automated triage

Benign / Suspicious / Malicious on every alert, with routing.

5

Deep investigation

Item exposure and device posture gathered automatically.

6

Analyst-gated response

Org-wide blocklisting only after a human approves.

7

Proactive hunting

Scheduled MCP-server audit across the estate.

8


Fleet script execution

Scheduled KOI script runs on Cortex-Agent endpoints.


ARCHITECTURE

How it flows

Telemetry pipeline




Alert-driven automation



Egress matters

KOI restricts its sensitive endpoints by source IP. A valid key can still return 403 from a shared cloud egress — run the integration on a Cortex engine whose egress IP KOI accepts. Event collection is unaffected.

Every alert, investigated before it reaches a human



Two investigations, run automatically

I Item investigation

What is this thing, and where is it?

- Catalog risk score, level and AI-generated risk summary
- Organization exposure — installs, publisher, signing status
- The endpoints it is installed on, and the users affected
- Governance state — is it already on the blocklist?
- Remediation and prior-approval history for the item

Runs inside triage, and standalone for an ad-hoc item lookup.


D Device investigation

What else is on the affected host?

- Every item installed on the device, with version and path
- Which of those sit at or above your risk threshold
- Marketplace, publisher and install location per item
- Remediations recorded against that host
- Turns a single alert into full host posture

Runs inside triage, or on any device id you hand it.

Action, gated on a human



0

org-wide blocks without an explicit analyst approval

Triage always invokes the response playbook with `auto_block = false`. On a Malicious verdict the run parks on the approval task; the blocklist command does not execute until a human approves. Verified on a live tenant.

Safe by default

The only state-changing command in the set is gated.

No redundant blocks

Items already on the blocklist short-circuit.

Full context at the gate

The approval shows risk, exposure and history.

26 commands across the KOI platform

7

Governance & policies

Policies, runtime hardening, findings, approvals, remediations

6

Allowlist & blocklist

Read and modify organization-wide allow and block lists

4

Devices & organization

Devices, per-device inventory, groups and users

4

Inventory

Items, item detail, search, and per-item endpoint exposure

3

Collector & diagnostics

Manual event pulls and fetch-state inspection

2

Koidex catalog

Catalog search and full item risk reports

Read-only by default — only the allow/block list and policy-status commands change tenant state.

Getting it running

1

Marketplace

- Search for KOI and install
- Integration, rules and dashboard land together
- Recommended for production tenants

2

Pack zip

- Build with demisto-sdk zip-packs
- Upload as a single object
- Best for dev/test and CI pipelines

3

Manual items

- Import individual playbooks and rules
- Import sub-playbooks before their callers
- Keep item names unchanged

Then prove it works

```
!koi-devices-list limit=5
```

```
dataset = koi_koi_raw | comp count() as events by source_log_type
```

Attach KOI - Alert Triage to a KOI alert and confirm the investigation summaries and verdict appear in the war room.

VALIDATION

Proven on a live tenant

11/11

read commands returned live data
— zero 403s once egress was
correct

500

items inventoried on a single host,
35 of them flagged high risk

3

playbook stages per alert — extract,
item and device investigation

0

blocklist writes fired before an
analyst approved

End-to-end run on a real alert

A vulnerable-dependency alert on a package whose catalog risk is High was escalated to **Malicious** — driven by real KOI catalog data, not the alert's own severity. The run then opened an analyst-gated block and parked on the approval task, with the blocklist command unexecuted.

Both investigations posted their summaries to the war room; the playbook completed cleanly.

WHERE TO GO NEXT

Install it, point it at KOI, attach the triage.

1

Customer guide

Install, configure, all 26 commands, playbooks, validation and troubleshooting — Word and PDF.

2

Playbooks reference

Architecture and configuration for both playbook suites.

3

Pack repository

Source for the integration, rules, dashboard and playbooks.