

KOI CONTENT PACK

Test Guide

Twelve tests that prove every part of the pack — with the tenant setup each one needs

For anyone comfortable in the Cortex XSIAM console. No content development required.

CONTENTS

What ships in the pack

A

KOI integration

Event collector plus 26 commands — devices, inventory, catalog risk, governance.

B

Parsing & modeling rules

Normalise KOI events and map them to the Cortex Data Model.

C

Alerts dashboard

Ready-made view of KOI alert activity.

D

Triage & investigation playbooks

7 playbooks: triage, item and device investigation, gated response, MCP hunt.

E

Script Runner playbooks

5 playbooks plus the KoiScanTracker automation — run KOI scripts fleet-wide.

F

XQL query library

72 saved queries: 26 threat hunts and 46 detections.

G

Test tooling

An event simulator, so you can test without waiting for real KOI activity.

Every one of these is covered by a test in this guide.

Every test is laid out the same way

PACK CONTENT

The content items that must be on the tenant for this test — playbooks, automations, rules, Lists. Useful if you upload the pack piece by piece rather than in one go.

TENANT SETUP

What you configure yourself: integration instances, an endpoint group, a script in Action Center, the configuration List.

STEPS + WHAT TO EXPECT

What to click, in order, in the Cortex XSIAM console — and what a passing test looks like. Anything to paste is shown in monospace.

Run them in order the first time

Tests 1 to 4 prove data is arriving and correct. Tests 5 to 8 prove the automation. Tests 9 and 10 prove fleet script execution — the only ones needing the Core REST API integration. Tests 11 and 12 are the query library and the simulator, and run any time. Nothing here needs a Cortex XDR integration: XSIAM is the XDR, and its XQL engine is built in.

Each test lists its own prerequisites — set those up first and the test runs straight through.

BEFORE YOU BEGIN

Everything the pack depends on, and where to set it up

What	Where in XSIAM	Needed for
KOI API key	KOI console → Settings → API Access	Tests 1-8, 11
KOI integration instance	Settings → Data Sources → Add → KOI	Tests 1-8, 11
An egress IP KOI accepts	Run the instance on a Cortex engine if needed	Tests 1-8, 11
Core REST API instance	Settings → Data Sources → Add → Core REST API	Tests 9-10 only
KOI script in Action Center	Action Center → Scripts Library → upload	Tests 9-10
An endpoint group	Tag the agents, then Endpoints → Endpoint Groups → dynamic	Tests 9-10
"Koi Script Runner" JSON List	Settings → Object Setup → Lists	Tests 9-10
Mail sender instance (optional)	Settings → Data Sources	Tests 9-10 notifications



The Core REST API instance is required for tests 9 and 10 only. The Script Runner reads endpoint groups through it, which is how it covers groups larger than 100 endpoints. Everything else in this guide runs without it.

Two ways in — only one collects events

Method	What lands	Events?
demisto-sdk upload	Everything — integration, playbooks, rules, dashboard	✓ yes
Pack zip upload	Content only. The collector and rules do not activate	✗ no

Ask whoever sent you the pack to install it

Installing with demisto-sdk is a developer task and needs a Standard XSIAM API key. If you are testing rather than deploying, have it installed for you and start at Test 1. Nothing else in this guide needs a developer.

Installed the pack, or uploading piece by piece?

- ✓ Installed the whole pack: every playbook, automation, rule and the Koi Script Runner List are already on the tenant. Nothing to upload.
- ✓ Uploading selectively: the PACK CONTENT column on each test is your upload checklist for that test.

TEST PLAN

Twelve tests, four groups

Data is arriving and correct

1 Connectivity

2 Event collection

3 Parsing & data model

4 Dashboard

The automation works

5 Alert triage

6 Item & device investigation

7 Gated response

8 MCP server hunt

Fleet script execution

prep What to put in place

9 Refresh the tracker

10 Scan due endpoints

Extras, any time

11 Threat-hunting library

12 Simulated test data



Only tests 9 and 10 need the Core REST API integration. Only tests 1-8 and 11 need a KOI API key. Test 12 needs neither and is the fastest way to see the pack working.

TEST 1

Connectivity and authorisation

PACK CONTENT

- KOI integration

TENANT SETUP

- A KOI API key (KOI console → Settings → API Access)
- Permission to add an integration instance in XSIAM

Steps

- 1 Settings → Data Sources → Add an instance → search KOI.
- 2 **Server URL:** `https://api.prod.koi.security/`
- 3 Paste the API key. Optionally set a Tenant Name label.
- 4 Click Test, then Save.
- 5 Open the CLI at the bottom of any incident and run:
- 6 `!koi-devices-list limit=5`

What to expect

- ✓ Test returns Success.
- ✓ The command returns a table of up to five devices.
- ✓ If you set a Tenant Name it will appear on collected events later.

TEST 2

Event collection

PACK CONTENT

- KOI integration
- Koi Parsing Rule

TENANT SETUP

- Test 1 passing
- Fetch events enabled on the KOI instance

Steps

- 1 Edit the KOI instance and tick Fetch events.
- 2 Choose which event types to collect, then Save.
- 3 Wait for two collection cycles.
- 4 Open Incident Response → Investigation → Query Builder and run:
- 5 `dataset = koi_koi_raw`
- 6 `| comp count() as rows by source_log_type`

What to expect

- ✓ Rows appear under Alerts, Audit, or both.
- ✓ Counts grow between cycles.

TEST 3

Parsing and the data model

PACK CONTENT

- Koi Parsing Rule
- Koi Modeling Rule

TENANT SETUP

- Test 2 passing — events present in koi_koi_raw

Steps

- 1 In Query Builder, count real alerts rather than rows:
- 2 `dataset = koi_koi_raw`
- 3 `| filter source_log_type = "Alerts"`
- 4 `| alter nid = json_extract_scalar(metadata, "$.notification_event_id")`
- 5 `| comp count() as rows, count_distinct(nid) as real_alerts`

What to expect

- ✓ Both numbers return.
- ✓ rows is larger than real_alerts — often much larger.
- ✓ Fields such as item_id, alert_hostname and risk_level are populated and directly queryable.

TEST 4

Dashboard

PACK CONTENT

- Koi Alerts Dashboard
- Koi Parsing Rule

TENANT SETUP

- Test 2 passing — events present

Steps

- 1 Dashboards & Reports → open the KOI Alerts Dashboard.
- 2 Read the alert totals, risk mix and top items.
- 3 Come back after another collection cycle and compare.

What to expect

- ✓ Widgets render with data.
- ✓ Alert counts are far lower than the raw row count from Test 3.
- ✓ Counts stay stable as fetches repeat — they do not climb on their own.

TEST 5

Alert triage

PACK CONTENT

- KOI IR - Alert Triage
- KOI IR - Extract Alert Context
- KOI IR - Investigate Item / Device
- incident field: KOI Notification ID

TENANT SETUP

- Test 1 passing — the KOI integration answers
- A KOI alert in XSIAM (a correlation rule over koi_koi_raw, or Test 12 for simulated data)

Steps

- 1 Open a KOI alert.
- 2 Attach the triage playbook, or run from the CLI:
- 3 `!setPlaybook playbookId="KOI IR - Alert Triage"`
- 4 Watch the Work Plan, then read the War Room.

What to expect

- ✓ Context is extracted: item, marketplace, risk, host.
- ✓ A verdict is posted: Malicious, Suspicious or Benign.
- ✓ Low risk auto-closes; critical and high escalate.
- ✓ Medium and pending stay open for a human.

TEST 6

Item and device investigation

PACK CONTENT

- KOI IR - Investigate Item
- KOI IR - Investigate Device
- KOI IR - Enrich Item

TENANT SETUP

- Test 1 passing
- An item id and its marketplace, and a device id — both visible on any KOI alert

Steps

- 1 Automation → Playbooks → KOI IR - Investigate Item → Run.
- 2 Supply item_id and marketplace.
- 3 Repeat with KOI IR - Investigate Device and a device id.

What to expect

- ✓ An investigation summary with catalog risk and an AI-written summary.
- ✓ Organisation exposure: installs, endpoints, affected users.
- ✓ Both governance counts — on blocklist AND on allowlist.
- ✓ For the device: everything installed on that host, and which of it is risky.

TEST 7

Gated response

PACK CONTENT

- KOI IR - Block and Remediate
- KOI IR - Investigate Item

TENANT SETUP

- Test 1 passing
- An item id and marketplace
- Permission to approve a task

Steps

- 1 Automation → Playbooks → KOI IR - Block and Remediate → Run.
- 2 Supply item_id and marketplace.
- 3 Open the pending approval task and read it.
- 4 Approve, or leave it pending.

What to expect

- ✓ The run PARKS on an approval task and waits.
- ✓ The approval shows the full investigation and both governance counts.
- ✓ Nothing reaches the blocklist until a human approves.
- ✓ An item already blocked short-circuits instead of being added twice.

TEST 8

MCP server hunt

PACK CONTENT

- KOI Hunting - MCP Server Audit

TENANT SETUP

- Test 1 passing — the KOI integration answers

Steps

- 1 Automation → Playbooks → KOI Hunting - MCP Server Audit → Run.
- 2 Optionally set risk_levels (default high, critical).
- 3 Read the War Room table.
- 4 To schedule it: Automation → Jobs → New Job, time-triggered.

What to expect

- ✓ A table of MCP servers at or above the risk threshold.
- ✓ An empty table is a valid result — it means nothing risky is installed.

Script Runner: everything this use case needs

PACK CONTENT — all of it

Five playbooks — import child before parent:

- 1 KOI Script Runner - Execute Endpoint Script
- 2 KOI Script Runner - Process Config Entry
- 3 KOI Script Runner - Refresh Entry
- 4 KOI Script Runner - Refresh Job
- 5 KOI Script Runner - Scan Job

Plus:

KoiScanTracker — an automation, not a playbook, and separate from the KOI deployment script opposite. Manual import takes one file: `dist/automation-KoiScanTracker.yml`

TENANT SETUP

- ✓ Core REST API integration instance — the Refresh job reads endpoint groups through it.
- ✓ The KOI deployment script in Action Center → Scripts Library. Download it from YOUR KOI console; it is KOI's script, not a Cortex one, and must take no parameters.
- ✓ An endpoint group — tag the agents, then build a dynamic group on that tag.

Do this once, before tests 9 and 10

- 1 Installed the whole pack? All of the above is already on the tenant — go straight to test 9.
- 2 Uploading selectively? Automation → Playbooks → Import for the five playbooks, in the order listed.
- 3 Automation → Scripts → Import → `dist/automation-KoiScanTracker.yml`
- 4 Settings → Object Setup → Lists → New List, type JSON, named exactly Koi Script Runner. Paste the block from Lists/README.md.
- 5 Edit two things in that List: `endpoint_groups` and `script.name`. One entry per OS; the rest has working defaults.

Script Runner — refresh the tracker

PACK CONTENT

- KOI Script Runner - Refresh Job
- KOI Script Runner - Refresh Entry
- KoiScanTracker (automation)
- List: Koi Script Runner

TENANT SETUP

- Everything on the previous slide is in place

Steps

- 1 Automation → Jobs → New Job, time-triggered, playbook KOI Script Runner - Refresh Job.
- 2 Enable the Job, then use Run now.

What to expect

- ✓ A tracker List appears under Lists, named as in your entry.
- ✓ It fills with one row per endpoint: an id and a last-scan value.
- ✓ Re-running adds new endpoints without disturbing existing rows.

Script Runner — scan due endpoints

PACK CONTENT

- KOI Script Runner - Scan Job
- KOI Script Runner - Process Config Entry
- KOI Script Runner - Execute Endpoint Script
- KoiScanTracker (our automation)

TENANT SETUP

- Test 9 passing — the tracker List has rows
- Connected agents in the group, matching the entry's endpoint_os

Steps

- 1 Automation → Jobs → New Job → time-triggered → playbook KOI Script Runner - Scan Job.
- 2 Enable the Job, then Run now.
- 3 Open the run, then check Action Center.
- 4 Run it a second time immediately.

What to expect

- ✓ The script is dispatched to due, connected endpoints.
- ✓ Those endpoints get a last-scan timestamp in the tracker.
- ✓ Offline, wrong-OS and unknown endpoints stay due and retry later.
- ✓ The second run does nothing — everything is inside the rescan interval.
- ✓ A SKIPPED entry means nothing was due; a STILL RUNNING entry means Action Center is finishing on its own. Both are normal and send no email.

TEST 11

Threat-hunting query library

PACK CONTENT

- none — the queries are files you paste

TENANT SETUP

- Test 2 passing — events present
- Nothing else — the cross-dataset hunts read xdr_data, which XSIAM holds natively

Steps

- 1 Open docs/xql in the pack you were sent.
- 2 Read QUERY_LIBRARY.md first — it explains the rules.
- 3 Paste a hunt into Query Builder, starting with H2.1.
- 4 Try a detection too, for example A1.

What to expect

- ✓ Each query runs and returns rows, or a clean empty result.
- ✓ Hunts surface findings KOI scored that nobody actioned.
- ✓ 26 hunts and 46 detections are included.

Simulated test data

PACK CONTENT

- KoiSimulateEvents (test automation)

TENANT SETUP

- Automation → Scripts → Import → TestTools/KoiSimulateEvents/automation-KoiSimulateEvents.yml (one file, code inside)
- No KOI API key and no Core REST API needed

Steps

- 1 Open the CLI on any incident.
- 2 Run:
- 3 `!KoiSimulateEvents alerts=40 duplicate_factor=8 audit=60`
- 4 Read the summary it prints — it tells you the answers in advance.
- 5 Query `kosisim_kosisim_raw` to check them.

What to expect

- ✓ 380 events land in `kosisim_kosisim_raw`.
- ✓ 40 distinct alerts delivered 8 times each, as real KOI does.
- ✓ The counts you measure match the summary exactly.

One line of evidence per test

- 1 Test returns Success; devices list returns rows
- 2 koi_koi_raw grows between collection cycles
- 3 Distinct alerts far below raw rows — and that is correct
- 4 Dashboard renders; counts stay stable across fetches
- 5 A verdict is reached; low risk auto-closes
- 6 Investigation shows catalog risk and both governance counts
- 7 Response parks on approval; blocklist untouched until approved
- 8 MCP audit returns servers at or above the threshold
- 9 Tracker List is created and fills with endpoints
- 10 Scan marks only due, connected endpoints; re-run is a no-op
- 11 Library queries run and return results
- 12 Simulated counts match the summary exactly

All twelve green — the pack is ready for production use.