

KOI CONTENT PACK • ACCEPTANCE TESTING

Test Guide

Nine tests that prove every part of the pack works on your tenant — and what a pass looks like for each.

9

tests

~45

minutes

1

tenant

0

risk to prod

PREREQUISITES

Before you start

Have these in place first — most failed tests trace back to a missing item here.

1

A KOI API key

KOI console → Settings → API Access. Needs the xt-Administrator role.

2

The pack installed

Integration, parsing/modeling rules, dashboard and playbooks on the tenant.

3

An egress path KOI accepts

KOI restricts sensitive endpoints by source IP. Cloud egress often returns 403.

4

A KOI alert to work with

At least one ingested alert with source koi, for tests 5 to 7.

5

For test 9 only

A script package in the Scripts Library, an endpoint group, and the Koi Script Runner List.

6

Nothing in production

Every test is read-only except the one you explicitly approve in test 7.

What you are going to prove

1

Connectivity & authorisation

The key works and your egress is accepted

2

Event collection

Alerts and audit logs land in koi_koi_raw

3

Command surface

The read commands return live tenant data

4

Dashboard & data model

Events are normalised and visualised

5

Alert triage end to end

Every alert is investigated and classified

6

The two investigations

Item and device context is gathered

7

Gated response

Nothing is blocked without a human

8

Scheduled hunt

Risky MCP servers surface without an alert

9

Script Runner job

KOI scripts run on Cortex-Agent endpoints

Run them in order — each test assumes the previous one passed.

TEST 1

Connectivity & authorisation

Steps

- 1 Settings → Integrations → add a KOI instance.
- 2 Server URL: `https://api.prod.koi.security/`
- 3 Paste the API key, then click Test.
- 4 `!koi-devices-list limit=5`
- 5 `!koi-users-list limit=1`

What to expect

- ✓ Test returns Success.
- ✓ A device table with up to five rows.
- ✓ koi-users-list returns a user record.
- ✓ No 403 Forbidden on either command.



koi-users-list is the egress probe. A valid key can still return 403 from a shared cloud egress — if it does, run the integration on a Cortex engine whose egress IP KOI accepts. Event collection is unaffected either way.

TEST 2

Event collection

Steps

- 1 On the instance, enable Fetch events and save.
- 2 Wait for one or two fetch cycles.
- 3 `dataset = koi_koi_raw`
- 4 `| comp count() as events by source_log_type`
- 5 Re-run the query after the next cycle.
- 6 `!koi-fetch-context-get`

What to expect

- ✓ Rows for the Alerts and Audit log types.
- ✓ Counts increase between the two runs.
- ✓ `koi-fetch-context-get` shows the cursor.



Event collection uses a different KOI endpoint from the command surface, so it keeps working even when the sensitive endpoints return 403. A passing test 2 alongside a failing test 1 points squarely at egress.

TEST 3

Command surface

Steps

- 1 Sweep the read commands in the playground:
- 2 `!koi-inventory-list limit=5`
- 3 `!koi-koidex-search query=<package>`
- 4 `!koi-blocklist-get`
- 5 `!koi-policy-list`
- 6 `!koi-findings-list`
- 7 Take a device id from test 1 and run:
- 8 `!koi-device-inventory-get device_id=<id>`

What to expect

- ✓ Each returns a table plus context under Koi.*
- ✓ No 403 on any of them.
- ✓ Lists everything installed on that host.
- ✓ 500 items on one host, 35 flagged high risk.



Twenty of the twenty-six commands are read-only. Only six change tenant state: the allowlist and blocklist add/remove pairs, koi-policy-status-update and koi-fetch-context-set. None of them run in this test.

TEST 4

Dashboard & data model

Steps

- 1 Dashboards & Reports → open Koi Alerts Dashboard.
- 2 Set the time range to cover your ingested data.
- 3 Confirm each widget renders.
- 4 Then check the modelling with XQL:
- 5 `dataset = koi_koi_raw | fields xdm.*`

What to expect

- ✓ Widgets render with data, not empty states.
- ✓ XDM fields are populated, not null.
- ✓ Alert activity matches what KOI shows.



An empty dashboard with a passing test 2 usually means the time range, not the data. Widen it before investigating the rules.

Alert triage, end to end

Steps

- 1 Open an ingested KOI alert (source koi).
- 2 Attach and run KOI - Alert Triage.
- 3 Watch the Work Plan to completion.
- 4 Read the war room from top to bottom.

What to expect

- ✓ Work Plan completes — runStatus: completed.
- ✓ An item investigation summary is posted.
- ✓ A device investigation summary is posted.
- ✓ A triage summary carries a verdict.
- ✓ Benign auto-closes, Suspicious stays open, Malicious raises severity.



Reference run: a vulnerable-dependency alert on a package whose KOI catalog risk is High escalated to Malicious — driven by the independent catalog data, not the alert's own severity. That corroboration is the point of the verdict logic.

TEST 6

The two investigations

Steps

- 1 In the test 5 run, open the item investigation summary.
- 2 Open the device investigation summary and its risky-items table.
- 3 Optionally run the device investigation on its own:
- 4 `KOI - Investigate Device, device_id=<id>`

What to expect

- ✓ Item: catalog risk and AI summary, org exposure, endpoints and users, blocklist state, remediation and approval history.
- ✓ Device: everything installed on the host, which of it is risky, host remediations.



Known cosmetic issue: a few item-investigation fields render with array brackets — ["high"] rather than high — because XSOAR resolves arrays differently inside a sub-playbook. The values are correct and the analyst-facing triage summary renders clean.

Response stays gated on a human

Steps

- 1 Let a Malicious verdict reach the response step, or run it directly:
- 2 `KOI - Block and Remediate`
- 3 `item_id=<id> marketplace=<mp> auto_block=false`
- 4 Do not approve yet.
- 5 Inspect the run state and the war room.
- 6 Then approve, and re-inspect.

What to expect

- ✓ The run parks — runStatus: waiting.
- ✓ koi-blocklist-items-add has NOT executed.
- ✓ The approval shows the full investigation.
- ✓ An already-blocklisted item short-circuits as Already blocked.
- ✓ Only then does the blocklist write fire.



This is the one test that must pass before you use the pack in production. Triage always calls the response playbook with `auto_block=false`, so a Malicious verdict can never block software on its own.

Scheduled hunt for risky MCP servers

Steps

- 1 Run KOI - MCP Server Audit manually first.
- 2 Leave risk_levels at its default:
- 3 `high,critical`
- 4 Review what it flags.
- 5 Then attach it to a time-triggered Job.

What to expect

- ✓ MCP-server inventory is enumerated.
- ✓ Items at or above the threshold are flagged and reported.
- ✓ Runs on your schedule with no alert needed.



The audit queries koi-inventory-list with view=mcp_servers. Confirm that view matches how your tenant categorises MCP servers — if it returns nothing, check the categorisation before assuming the estate is clean.

Script Runner job

Steps

- 1 Confirm the prerequisites exist and match by name.
- 2 Create the JSON List named exactly:
- 3 `Koi Script Runner`
- 4 Automation → Jobs → Run now.
- 5 Open the run, then open Action Center.

What to expect

- ✓ Work Plan is green.
- ✓ ScriptResult ok:true plus an action_id.
- ✓ SKIPPED info entries where no connected endpoint matches the OS scope — and no failure email.
- ✓ Action Center shows COMPLETED_SUCCESSFULLY.



Everything binds by name at run time: the script package in the Scripts Library, the endpoint groups, the List and the mail instance. Rename any of them without updating the List and the next run fails or skips — the war room records which reference broke.

IF SOMETHING FAILS

Symptom → cause → fix

403 on inventory, koidex or users

Source IP not accepted by KOI's edge

Run the integration on a Cortex engine with an allowlisted egress IP

Test passes but no events arrive

Fetch events not enabled, or no new KOI activity

Enable it, wait two cycles, check koi-fetch-context-get

Triage cannot find the item

The alert maps the KOI payload elsewhere

Adjust alert_description on KOI - Extract Alert Context

Enrichment empty, playbook still green

Best-effort by design — KOI was unreachable

Verify with !koi-users-list limit=1

Job entry reports SKIPPED

No connected, unisolated endpoint matches the OS scope

Check group membership and endpoint_os in the List

Script not found at dispatch

Library name does not match script.name exactly

Fix the name, or pin script.uuid to be rename-proof

The first row accounts for most real-world failures — check egress before anything else.

One line of evidence per test

1

Test returns Success and koi-devices-list returns rows

2

koi_koi_raw event counts grow between fetch cycles

3

Read commands return tenant data with no 403

4

Dashboard widgets render and XDM fields are populated

5

Triage completes with a verdict and both summaries

6

Item and device investigations carry real KOI data

7

Run parks on approval; blocklist write did not execute

8

MCP audit flags items at or above the threshold

9

Job returns ScriptResult ok with an action_id

All nine green → the pack is ready for production use.